

# brain tumor with Crop image

January 10, 2023

## 1 Brain\_Tumor with crop image

```
[1]: !pip install imutils
```

Requirement already satisfied: imutils in  
/home/tanzim/anaconda3/lib/python3.9/site-packages (0.5.4)

```
[2]: import numpy as np
from tqdm import tqdm
import cv2
import os
import imutils
```

```
[4]: def crop_img(img):
    """
    Finds the extreme points on the image and crops the rectangular out of
    them
    """
    gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
    gray = cv2.GaussianBlur(gray, (3, 3), 0)

    # threshold the image, then perform a series of erosions +
    # dilations to remove any small regions of noise
    thresh = cv2.threshold(gray, 45, 255, cv2.THRESH_BINARY)[1]
    thresh = cv2.erode(thresh, None, iterations=2)
    thresh = cv2.dilate(thresh, None, iterations=2)

    # find contours in thresholded image, then grab the largest one
    cnts = cv2.findContours(thresh.copy(), cv2.RETR_EXTERNAL, cv2.
    CHAIN_APPROX_SIMPLE)
    cnts = imutils.grab_contours(cnts)
    c = max(cnts, key=cv2.contourArea)

    # find the extreme points
    extLeft = tuple(c[c[:, :, 0].argmin()][0])
    extRight = tuple(c[c[:, :, 0].argmax()][0])
    extTop = tuple(c[c[:, :, 1].argmin()][0])
```

```

        extBot = tuple(c[c[:, :, 1].argmax()][0])
        ADD_PIXELS = 0
        new_img = img[extTop[1]-ADD_PIXELS:extBot[1]+ADD_PIXELS,
        ↪extLeft[0]-ADD_PIXELS:extRight[0]+ADD_PIXELS].copy()

        return new_img

if __name__ == "__main__":
    training = "Training"
    testing = "Testing"
    training_dir = os.listdir(training)
    testing_dir = os.listdir(testing)
    IMG_SIZE = 256

    for dir in training_dir:
        save_path = 'cleaned/Training/' + dir
        path = os.path.join(training, dir)
        image_dir = os.listdir(path)
        for img in image_dir:
            image = cv2.imread(os.path.join(path, img))
            new_img = crop_img(image)
            new_img = cv2.resize(new_img, (IMG_SIZE, IMG_SIZE))
            if not os.path.exists(save_path):
                os.makedirs(save_path)
            cv2.imwrite(save_path+'/'+img, new_img)

    for dir in testing_dir:
        save_path = 'cleaned/Testing/' + dir
        path = os.path.join(testing, dir)
        image_dir = os.listdir(path)
        for img in image_dir:
            image = cv2.imread(os.path.join(path, img))
            new_img = crop_img(image)
            new_img = cv2.resize(new_img, (IMG_SIZE, IMG_SIZE))
            if not os.path.exists(save_path):
                os.makedirs(save_path)
            cv2.imwrite(save_path+'/'+img, new_img)

```

```

[5]: # Necessary imports
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,
    ↪Dropout
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam, SGD

```

```

from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau,
↳ModelCheckpoint

import os
import cv2
import matplotlib.pyplot as plt
import seaborn as sns

```

2023-01-09 23:56:07.656752: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```


[6]: data_dir = ('cleaned/Training')
categories = ['glioma', 'meningioma', 'notumor', 'pituitary']
for i in categories:
    path = os.path.join(data_dir, i)
    for img in os.listdir(path):
        img_array = cv2.imread(os.path.join(path,img))

```

```

[7]: plt.imshow(img_array);

```







```
[8]: # The image shape.  
img_array.shape
```

```
[8]: (256, 256, 3)
```

## 2 Plotting a image of each brain tumor type

```
[9]: plt.figure(figsize=(20, 16))  
  
images_path = ['/glioma/Tr-glTr_0000.jpg', '/meningioma/Tr-meTr_0000.jpg', '/  
↳notumor/Tr-noTr_0000.jpg', '/pituitary/Tr-piTr_0000.jpg']  
  
for i in range(4):  
    ax = plt.subplot(2, 2, i + 1)  
    img = cv2.imread(data_dir + images_path[i])  
    img = cv2.resize(img, (250, 250))  
    plt.imshow(img)  
    plt.title(categories[i])
```



### 3 Modelling: CNN

```
[10]: model1 = Sequential()
      # Convolutional layer 1
      model1.add(Conv2D(32,(3,3), input_shape=(64, 64, 1), activation='relu'))
      model1.add(BatchNormalization())
      model1.add(MaxPooling2D(pool_size=(2,2)))

      # Convolutional layer 2
      model1.add(Conv2D(32,(3,3), activation='relu'))
      model1.add(BatchNormalization())
      model1.add(MaxPooling2D(pool_size=(2,2)))

      model1.add(Flatten())
```

```

# Neural network

model1.add(Dense(units= 252, activation='relu'))
model1.add(Dropout(0.2))
model1.add(Dense(units=252, activation='relu'))
model1.add(Dropout(0.2))
model1.add(Dense(units=4, activation='softmax'))
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001, decay=0.0001,
clipvalue=0.5)
model1.compile(optimizer=optimizer, loss='categorical_crossentropy',
               metrics= ['categorical_accuracy'])

# using the ImageDataGenerator to prepare the images (Resize, nomalize, etc)

generator_train = ImageDataGenerator(rescale=1./255,
                                     featurewise_center=False,
                                     samplewise_center=False,
                                     featurewise_std_normalization=False,
                                     samplewise_std_normalization=False,
                                     zca_whitening=False,
                                     rotation_range=0,
                                     zoom_range = 0,
                                     width_shift_range=0,
                                     height_shift_range=0,
                                     horizontal_flip=True,
                                     vertical_flip=False)

generator_test = ImageDataGenerator(rescale=1./255,
                                    featurewise_center=False,
                                    samplewise_center=False,
                                    featurewise_std_normalization=False,
                                    samplewise_std_normalization=False,
                                    zca_whitening=False,
                                    rotation_range=0,
                                    zoom_range = 0,
                                    width_shift_range=0,
                                    height_shift_range=0,
                                    horizontal_flip=True,
                                    vertical_flip=False)

```

2023-01-09 23:58:32.623119: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2023-01-09 23:58:32.623679: I

tensorflow/core/common\_runtime/process\_util.cc:146] Creating new thread pool with default inter op setting: 2. Tune using inter\_op\_parallelism\_threads for best performance.

[11]: *# Creating the train and test data.*

```
train = generator_train.flow_from_directory('cleaned/Training',  
    ↪target_size=(64,64),  
                                           batch_size=32, class_mode=  
    ↪"categorical", color_mode='grayscale')  
  
test = generator_test.flow_from_directory('cleaned/Testing',  
    ↪target_size=(64,64),  
                                           batch_size=32, class_mode=  
    ↪"categorical", color_mode='grayscale')
```

Found 5712 images belonging to 4 classes.

Found 1311 images belonging to 4 classes.

[12]: *# Creating callbacks for the model.*

```
# If the model doesn't continue to improve (loss), the training will stop.  
  
# Stop training if loss doesn't keep decreasing.  
model1_es = EarlyStopping(monitor = 'loss', min_delta = 1e-11, patience = 12,  
    ↪verbose = 1)  
model1_rlr = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.2, patience =  
    ↪6, verbose = 1)  
  
# Automatically saves the best weights of the model, based on best val_accuracy  
model1_mcp = ModelCheckpoint(filepath = 'model1_weights.h5', monitor =  
    ↪'val_categorical_accuracy',  
                             save_best_only = True, verbose = 1)  
  
# Fiting the model.  
history1 = model1.fit(train, steps_per_epoch=5712//32, epochs=100,  
    ↪validation_data=test, validation_steps= 1311//32,  
                      callbacks=[model1_es, model1_rlr, model1_mcp])
```

Epoch 1/100

178/178 [=====] - ETA: 0s - loss: 0.8936 -  
categorical\_accuracy: 0.7018

Epoch 1: val\_categorical\_accuracy improved from -inf to 0.51172, saving model to  
model1\_weights.h5

178/178 [=====] - 36s 199ms/step - loss: 0.8936 -

```

categorical_accuracy: 0.7018 - val_loss: 4.4978 - val_categorical_accuracy:
0.5117 - lr: 0.0010
Epoch 2/100
178/178 [=====] - ETA: 0s - loss: 0.4307 -
categorical_accuracy: 0.8433
Epoch 2: val_categorical_accuracy did not improve from 0.51172
178/178 [=====] - 49s 277ms/step - loss: 0.4307 -
categorical_accuracy: 0.8433 - val_loss: 6.3103 - val_categorical_accuracy:
0.3789 - lr: 0.0010
Epoch 3/100
178/178 [=====] - ETA: 0s - loss: 0.3027 -
categorical_accuracy: 0.8875
Epoch 3: val_categorical_accuracy improved from 0.51172 to 0.56172, saving model
to model1_weights.h5
178/178 [=====] - 49s 276ms/step - loss: 0.3027 -
categorical_accuracy: 0.8875 - val_loss: 2.7020 - val_categorical_accuracy:
0.5617 - lr: 0.0010
Epoch 4/100
178/178 [=====] - ETA: 0s - loss: 0.2275 -
categorical_accuracy: 0.9208
Epoch 4: val_categorical_accuracy improved from 0.56172 to 0.77578, saving model
to model1_weights.h5
178/178 [=====] - 48s 268ms/step - loss: 0.2275 -
categorical_accuracy: 0.9208 - val_loss: 0.7948 - val_categorical_accuracy:
0.7758 - lr: 0.0010
Epoch 5/100
178/178 [=====] - ETA: 0s - loss: 0.1820 -
categorical_accuracy: 0.9371
Epoch 5: val_categorical_accuracy improved from 0.77578 to 0.89766, saving model
to model1_weights.h5
178/178 [=====] - 45s 252ms/step - loss: 0.1820 -
categorical_accuracy: 0.9371 - val_loss: 0.2957 - val_categorical_accuracy:
0.8977 - lr: 0.0010
Epoch 6/100
178/178 [=====] - ETA: 0s - loss: 0.1258 -
categorical_accuracy: 0.9583
Epoch 6: val_categorical_accuracy did not improve from 0.89766
178/178 [=====] - 46s 261ms/step - loss: 0.1258 -
categorical_accuracy: 0.9583 - val_loss: 0.5858 - val_categorical_accuracy:
0.7961 - lr: 0.0010
Epoch 7/100
178/178 [=====] - ETA: 0s - loss: 0.1161 -
categorical_accuracy: 0.9563
Epoch 7: val_categorical_accuracy improved from 0.89766 to 0.95000, saving model
to model1_weights.h5
178/178 [=====] - 42s 234ms/step - loss: 0.1161 -
categorical_accuracy: 0.9563 - val_loss: 0.1326 - val_categorical_accuracy:
0.9500 - lr: 0.0010

```



Epoch 8/100  
178/178 [=====] - ETA: 0s - loss: 0.0940 -  
categorical\_accuracy: 0.9699  
Epoch 8: val\_categorical\_accuracy did not improve from 0.95000  
178/178 [=====] - 45s 253ms/step - loss: 0.0940 -  
categorical\_accuracy: 0.9699 - val\_loss: 0.2149 - val\_categorical\_accuracy:  
0.9312 - lr: 0.0010  
Epoch 9/100  
178/178 [=====] - ETA: 0s - loss: 0.0786 -  
categorical\_accuracy: 0.9754  
Epoch 9: val\_categorical\_accuracy improved from 0.95000 to 0.95078, saving model  
to model1\_weights.h5  
178/178 [=====] - 41s 228ms/step - loss: 0.0786 -  
categorical\_accuracy: 0.9754 - val\_loss: 0.1498 - val\_categorical\_accuracy:  
0.9508 - lr: 0.0010  
Epoch 10/100  
178/178 [=====] - ETA: 0s - loss: 0.0634 -  
categorical\_accuracy: 0.9810  
Epoch 10: val\_categorical\_accuracy improved from 0.95078 to 0.95547, saving  
model to model1\_weights.h5  
178/178 [=====] - 47s 262ms/step - loss: 0.0634 -  
categorical\_accuracy: 0.9810 - val\_loss: 0.1220 - val\_categorical\_accuracy:  
0.9555 - lr: 0.0010  
Epoch 11/100  
178/178 [=====] - ETA: 0s - loss: 0.0684 -  
categorical\_accuracy: 0.9778  
Epoch 11: val\_categorical\_accuracy did not improve from 0.95547  
178/178 [=====] - 47s 265ms/step - loss: 0.0684 -  
categorical\_accuracy: 0.9778 - val\_loss: 0.2696 - val\_categorical\_accuracy:  
0.9086 - lr: 0.0010  
Epoch 12/100  
178/178 [=====] - ETA: 0s - loss: 0.0508 -  
categorical\_accuracy: 0.9812  
Epoch 12: val\_categorical\_accuracy improved from 0.95547 to 0.97266, saving  
model to model1\_weights.h5  
178/178 [=====] - 42s 235ms/step - loss: 0.0508 -  
categorical\_accuracy: 0.9812 - val\_loss: 0.0784 - val\_categorical\_accuracy:  
0.9727 - lr: 0.0010  
Epoch 13/100  
178/178 [=====] - ETA: 0s - loss: 0.0579 -  
categorical\_accuracy: 0.9798  
Epoch 13: val\_categorical\_accuracy did not improve from 0.97266  
178/178 [=====] - 46s 261ms/step - loss: 0.0579 -  
categorical\_accuracy: 0.9798 - val\_loss: 0.2044 - val\_categorical\_accuracy:  
0.9375 - lr: 0.0010  
Epoch 14/100  
178/178 [=====] - ETA: 0s - loss: 0.0365 -  
categorical\_accuracy: 0.9877

Epoch 14: val\_categorical\_accuracy improved from 0.97266 to 0.97422, saving model to model1\_weights.h5  
178/178 [=====] - 46s 257ms/step - loss: 0.0365 - categorical\_accuracy: 0.9877 - val\_loss: 0.1115 - val\_categorical\_accuracy: 0.9742 - lr: 0.0010

Epoch 15/100  
178/178 [=====] - ETA: 0s - loss: 0.0505 - categorical\_accuracy: 0.9850

Epoch 15: val\_categorical\_accuracy did not improve from 0.97422  
178/178 [=====] - 46s 256ms/step - loss: 0.0505 - categorical\_accuracy: 0.9850 - val\_loss: 0.3364 - val\_categorical\_accuracy: 0.9172 - lr: 0.0010

Epoch 16/100  
178/178 [=====] - ETA: 0s - loss: 0.0505 - categorical\_accuracy: 0.9852

Epoch 16: val\_categorical\_accuracy did not improve from 0.97422  
178/178 [=====] - 44s 249ms/step - loss: 0.0505 - categorical\_accuracy: 0.9852 - val\_loss: 0.2937 - val\_categorical\_accuracy: 0.9359 - lr: 0.0010

Epoch 17/100  
178/178 [=====] - ETA: 0s - loss: 0.0295 - categorical\_accuracy: 0.9907

Epoch 17: val\_categorical\_accuracy did not improve from 0.97422  
178/178 [=====] - 42s 236ms/step - loss: 0.0295 - categorical\_accuracy: 0.9907 - val\_loss: 0.1335 - val\_categorical\_accuracy: 0.9680 - lr: 0.0010

Epoch 18/100  
178/178 [=====] - ETA: 0s - loss: 0.0347 - categorical\_accuracy: 0.9898

Epoch 18: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 18: val\_categorical\_accuracy did not improve from 0.97422  
178/178 [=====] - 46s 261ms/step - loss: 0.0347 - categorical\_accuracy: 0.9898 - val\_loss: 0.1507 - val\_categorical\_accuracy: 0.9625 - lr: 0.0010

Epoch 19/100  
178/178 [=====] - ETA: 0s - loss: 0.0185 - categorical\_accuracy: 0.9933

Epoch 19: val\_categorical\_accuracy improved from 0.97422 to 0.98125, saving model to model1\_weights.h5  
178/178 [=====] - 72s 406ms/step - loss: 0.0185 - categorical\_accuracy: 0.9933 - val\_loss: 0.0919 - val\_categorical\_accuracy: 0.9812 - lr: 2.0000e-04

Epoch 20/100  
178/178 [=====] - ETA: 0s - loss: 0.0103 - categorical\_accuracy: 0.9965

Epoch 20: val\_categorical\_accuracy did not improve from 0.98125  
178/178 [=====] - 84s 474ms/step - loss: 0.0103 -

```

categorical_accuracy: 0.9965 - val_loss: 0.0881 - val_categorical_accuracy:
0.9781 - lr: 2.0000e-04
Epoch 21/100
178/178 [=====] - ETA: 0s - loss: 0.0109 -
categorical_accuracy: 0.9974
Epoch 21: val_categorical_accuracy did not improve from 0.98125
178/178 [=====] - 69s 388ms/step - loss: 0.0109 -
categorical_accuracy: 0.9974 - val_loss: 0.0989 - val_categorical_accuracy:
0.9758 - lr: 2.0000e-04
Epoch 22/100
178/178 [=====] - ETA: 0s - loss: 0.0073 -
categorical_accuracy: 0.9970
Epoch 22: val_categorical_accuracy did not improve from 0.98125
178/178 [=====] - 82s 459ms/step - loss: 0.0073 -
categorical_accuracy: 0.9970 - val_loss: 0.0939 - val_categorical_accuracy:
0.9805 - lr: 2.0000e-04
Epoch 23/100
178/178 [=====] - ETA: 0s - loss: 0.0090 -
categorical_accuracy: 0.9970
Epoch 23: val_categorical_accuracy did not improve from 0.98125
178/178 [=====] - 82s 463ms/step - loss: 0.0090 -
categorical_accuracy: 0.9970 - val_loss: 0.0869 - val_categorical_accuracy:
0.9789 - lr: 2.0000e-04
Epoch 24/100
178/178 [=====] - ETA: 0s - loss: 0.0110 -
categorical_accuracy: 0.9968
Epoch 24: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 24: val_categorical_accuracy improved from 0.98125 to 0.98281, saving
model to model1_weights.h5
178/178 [=====] - 93s 522ms/step - loss: 0.0110 -
categorical_accuracy: 0.9968 - val_loss: 0.1023 - val_categorical_accuracy:
0.9828 - lr: 2.0000e-04
Epoch 25/100
178/178 [=====] - ETA: 0s - loss: 0.0078 -
categorical_accuracy: 0.9975
Epoch 25: val_categorical_accuracy did not improve from 0.98281
178/178 [=====] - 92s 517ms/step - loss: 0.0078 -
categorical_accuracy: 0.9975 - val_loss: 0.0810 - val_categorical_accuracy:
0.9828 - lr: 4.0000e-05
Epoch 26/100
178/178 [=====] - ETA: 0s - loss: 0.0062 -
categorical_accuracy: 0.9979
Epoch 26: val_categorical_accuracy did not improve from 0.98281
178/178 [=====] - 93s 523ms/step - loss: 0.0062 -
categorical_accuracy: 0.9979 - val_loss: 0.0887 - val_categorical_accuracy:
0.9812 - lr: 4.0000e-05
Epoch 27/100

```

```

178/178 [=====] - ETA: 0s - loss: 0.0075 -
categorical_accuracy: 0.9975
Epoch 27: val_categorical_accuracy did not improve from 0.98281
178/178 [=====] - 92s 519ms/step - loss: 0.0075 -
categorical_accuracy: 0.9975 - val_loss: 0.0904 - val_categorical_accuracy:
0.9828 - lr: 4.0000e-05
Epoch 28/100
178/178 [=====] - ETA: 0s - loss: 0.0043 -
categorical_accuracy: 0.9988
Epoch 28: val_categorical_accuracy did not improve from 0.98281
178/178 [=====] - 93s 521ms/step - loss: 0.0043 -
categorical_accuracy: 0.9988 - val_loss: 0.0924 - val_categorical_accuracy:
0.9805 - lr: 4.0000e-05
Epoch 29/100
178/178 [=====] - ETA: 0s - loss: 0.0035 -
categorical_accuracy: 0.9991
Epoch 29: val_categorical_accuracy improved from 0.98281 to 0.98359, saving
model to model1_weights.h5
178/178 [=====] - 91s 509ms/step - loss: 0.0035 -
categorical_accuracy: 0.9991 - val_loss: 0.0882 - val_categorical_accuracy:
0.9836 - lr: 4.0000e-05
Epoch 30/100
178/178 [=====] - ETA: 0s - loss: 0.0046 -
categorical_accuracy: 0.9989
Epoch 30: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.

Epoch 30: val_categorical_accuracy did not improve from 0.98359
178/178 [=====] - 93s 523ms/step - loss: 0.0046 -
categorical_accuracy: 0.9989 - val_loss: 0.0986 - val_categorical_accuracy:
0.9773 - lr: 4.0000e-05
Epoch 31/100
178/178 [=====] - ETA: 0s - loss: 0.0058 -
categorical_accuracy: 0.9979
Epoch 31: val_categorical_accuracy did not improve from 0.98359
178/178 [=====] - 89s 500ms/step - loss: 0.0058 -
categorical_accuracy: 0.9979 - val_loss: 0.0921 - val_categorical_accuracy:
0.9820 - lr: 8.0000e-06
Epoch 32/100
178/178 [=====] - ETA: 0s - loss: 0.0057 -
categorical_accuracy: 0.9979
Epoch 32: val_categorical_accuracy did not improve from 0.98359
178/178 [=====] - 52s 293ms/step - loss: 0.0057 -
categorical_accuracy: 0.9979 - val_loss: 0.0991 - val_categorical_accuracy:
0.9789 - lr: 8.0000e-06
Epoch 33/100
178/178 [=====] - ETA: 0s - loss: 0.0048 -
categorical_accuracy: 0.9981
Epoch 33: val_categorical_accuracy did not improve from 0.98359

```

178/178 [=====] - 53s 297ms/step - loss: 0.0048 -  
categorical\_accuracy: 0.9981 - val\_loss: 0.0959 - val\_categorical\_accuracy:  
0.9789 - lr: 8.0000e-06  
Epoch 34/100  
178/178 [=====] - ETA: 0s - loss: 0.0064 -  
categorical\_accuracy: 0.9979  
Epoch 34: val\_categorical\_accuracy did not improve from 0.98359  
178/178 [=====] - 44s 246ms/step - loss: 0.0064 -  
categorical\_accuracy: 0.9979 - val\_loss: 0.0909 - val\_categorical\_accuracy:  
0.9805 - lr: 8.0000e-06  
Epoch 35/100  
178/178 [=====] - ETA: 0s - loss: 0.0054 -  
categorical\_accuracy: 0.9982  
Epoch 35: val\_categorical\_accuracy did not improve from 0.98359  
178/178 [=====] - 48s 269ms/step - loss: 0.0054 -  
categorical\_accuracy: 0.9982 - val\_loss: 0.0995 - val\_categorical\_accuracy:  
0.9820 - lr: 8.0000e-06  
Epoch 36/100  
178/178 [=====] - ETA: 0s - loss: 0.0061 -  
categorical\_accuracy: 0.9979  
Epoch 36: val\_categorical\_accuracy did not improve from 0.98359  
178/178 [=====] - 45s 255ms/step - loss: 0.0061 -  
categorical\_accuracy: 0.9979 - val\_loss: 0.0778 - val\_categorical\_accuracy:  
0.9820 - lr: 8.0000e-06  
Epoch 37/100  
178/178 [=====] - ETA: 0s - loss: 0.0050 -  
categorical\_accuracy: 0.9984  
Epoch 37: val\_categorical\_accuracy did not improve from 0.98359  
178/178 [=====] - 47s 266ms/step - loss: 0.0050 -  
categorical\_accuracy: 0.9984 - val\_loss: 0.0812 - val\_categorical\_accuracy:  
0.9820 - lr: 8.0000e-06  
Epoch 38/100  
178/178 [=====] - ETA: 0s - loss: 0.0057 -  
categorical\_accuracy: 0.9989  
Epoch 38: val\_categorical\_accuracy did not improve from 0.98359  
178/178 [=====] - 62s 347ms/step - loss: 0.0057 -  
categorical\_accuracy: 0.9989 - val\_loss: 0.0901 - val\_categorical\_accuracy:  
0.9820 - lr: 8.0000e-06  
Epoch 39/100  
178/178 [=====] - ETA: 0s - loss: 0.0040 -  
categorical\_accuracy: 0.9989  
Epoch 39: val\_categorical\_accuracy did not improve from 0.98359  
178/178 [=====] - 76s 427ms/step - loss: 0.0040 -  
categorical\_accuracy: 0.9989 - val\_loss: 0.1036 - val\_categorical\_accuracy:  
0.9820 - lr: 8.0000e-06  
Epoch 40/100  
178/178 [=====] - ETA: 0s - loss: 0.0049 -  
categorical\_accuracy: 0.9982

Epoch 40: val\_categorical\_accuracy did not improve from 0.98359  
 178/178 [=====] - 73s 412ms/step - loss: 0.0049 -  
 categorical\_accuracy: 0.9982 - val\_loss: 0.0883 - val\_categorical\_accuracy:  
 0.9805 - lr: 8.0000e-06  
 Epoch 41/100  
 178/178 [=====] - ETA: 0s - loss: 0.0057 -  
 categorical\_accuracy: 0.9981  
 Epoch 41: val\_categorical\_accuracy did not improve from 0.98359  
 178/178 [=====] - 63s 356ms/step - loss: 0.0057 -  
 categorical\_accuracy: 0.9981 - val\_loss: 0.0917 - val\_categorical\_accuracy:  
 0.9820 - lr: 8.0000e-06  
 Epoch 41: early stopping

[13]: model1.evaluate(test)

41/41 [=====] - 3s 72ms/step - loss: 0.0811 -  
 categorical\_accuracy: 0.9825

[13]: [0.08108488470315933, 0.9824561476707458]

[14]: model1.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	320
batch_normalization (Batch Normalization)	(None, 62, 62, 32)	128
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 29, 29, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 252)	1580796
dropout (Dropout)	(None, 252)	0

dense_1 (Dense)	(None, 252)	63756
dropout_1 (Dropout)	(None, 252)	0
dense_2 (Dense)	(None, 4)	1012

```
=====
Total params: 1,655,388
Trainable params: 1,655,260
Non-trainable params: 128
-----
```

```
[ ]:
```